

On the Effectiveness of Type-based Control Flow Integrity

Reza Mirzazade farkhani, Saman Jafari, Sajjad Arshad,
William Robertson, Engin Kirda, Hamed Okhravi



Northeastern University
College of Computer and Information Science

NEU SecLab





Outline




- Control Flow Hijacking
- Control Flow Integrity (CFI)
- Runtime Type Checking (RTC)
- Reuse Attack Protector (RAP)
- Typed ROP (TROP)
- PoC Exploit for Nginx
- Evaluation
- Conclusion




Control Flow Hijacking

- **Memory Corruption** can lead to violation of **Control Flow Graph (CFG)**

Forward Edges (calls, jumps)



Backward Edges (return addresses)





Control Flow Integrity (CFI)

- CFI prevents control flow hijacking by enforcing CFG at runtime
- CFG is usually generated *statically* using **Points-to Analysis**
 - **DSA**: Data Structure Analysis
 - **SVF**: Static Value-Flow Analysis
- Constructing **Sound** and **Precise** CFGs is **undecidable** and **impractical**



Runtime Type Checking (RTC)



- **Runtime Type Checking (RTC)** generates the CFG based on **Type Signature**
- RTC matches the type signature of each indirect control transfer with its target.
- Forward edge
 - The type of function pointer and the target are checked at each control transfer.
- Backward edge
 - The type of callee is checked during the function epilogue.
- Implementations
 - TypeArmor, KCFI (Kernel CFI), MCFI (Modular CFI), LLVM-CFI



Reuse Attack Protector (RAP)



**RAP™ is here. Public demo in 4.5 test patch
and commercially available today!**


April 28, 2016

**RAP™ Demonstrates World-First
Fully CFI-Hardened OS Kernel**

Type-based, high-performance, high-security, forward/backward-edge CFI
February 6, 2017



Reuse Attack Protector (RAP)





Reuse Attack Protector (RAP)

typedef	void	(*fptr) (int	,	long)	hash: 5d769299
	void	func1 (int	i,	long	j)	hash: 5d769299
	void	func2 (int	m,	long	n)	hash: 5d769299

Return Type Parameters Type

Type
Collision



Sample Vulnerable Program



```
1 typedef void (*FunctionPointer)(void);
2 int flag = 0;
3 char *cmd;
4 void valid_target1(void) {
5     printf("Valid Target 1\n");
6 }
7 void valid_target2(void) {
8     printf("Valid Target 2\n");
9 }
10
11 int final_target(char *cmd) {
12     system(cmd);
13 }
14 int linker_func(void) {
15     if (flag == 1)
16         final_target(cmd);
17 }
18 void invalid_target(void) {
19     linker_func();
20 }
21 void vulnerable(char * input) {
22     FunctionPointer corruptible_fptr;
23     char buf[20];
24     if (strcmp(input, "1") == 0)
25         corruptible_fptr = &valid_target1;
26     else
27         corruptible_fptr = &valid_target2;
28     printf(input);
29     strcpy(buf, input);
30     corruptible_fptr();
31 }
```

invalid_target()

hash2
...
call linker_func()

linker_func()

hash3
...
call final_target()

final_target()

hash4
...
call execve()

vulnerable()

hash1
...
call (*fptr)

valid_target1()

hash2
...

valid_target2()

hash2
...



Research Questions



- Can RTC be practically bypassed using **type collisions**?
- Are there enough intermediate functions with satisfiable constraints in real-world applications?
- How prevalent are these constructs in real-world applications?




Threat Model



- The attacker has arbitrary **read** and **write** primitives to the memory
- The application contains **one strong** or **multiple limited** memory corruption vulnerability
- DEP and ASLR are enabled
- RAP is in place

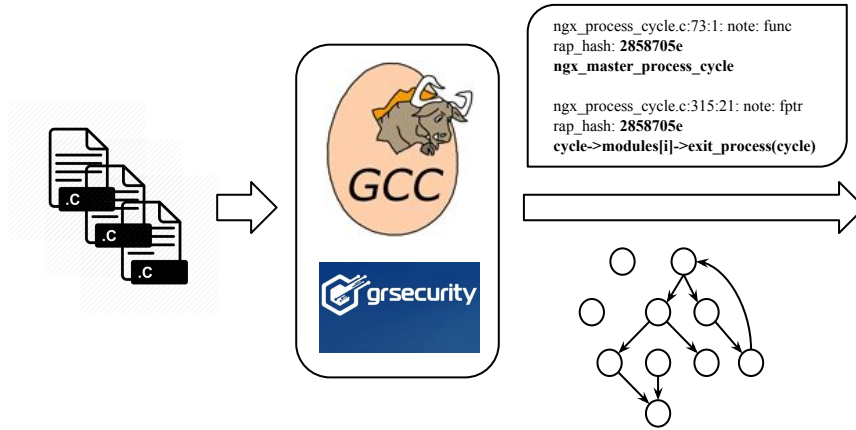


Typed ROP (TROP)



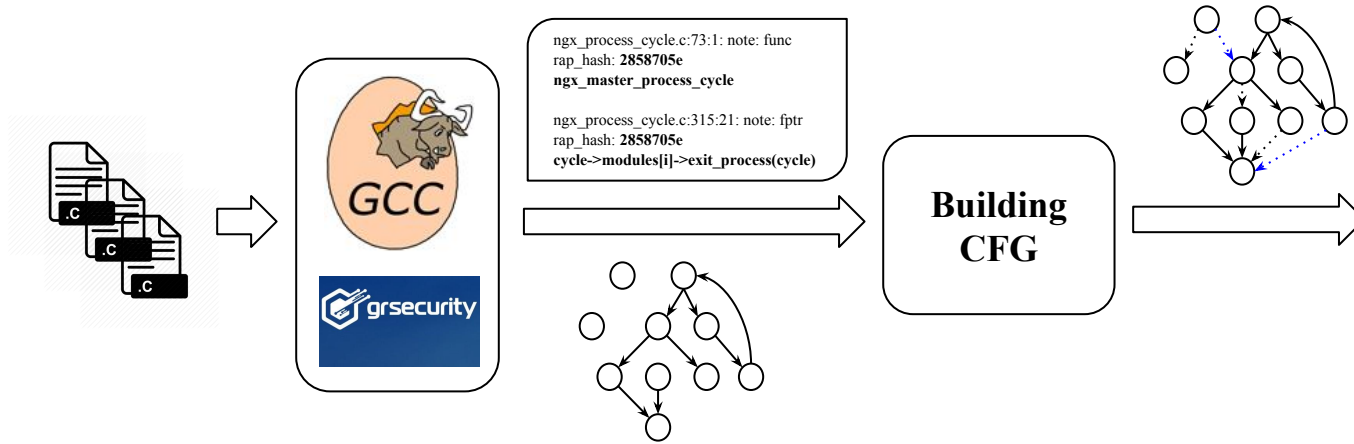


Typed ROP (TROP)



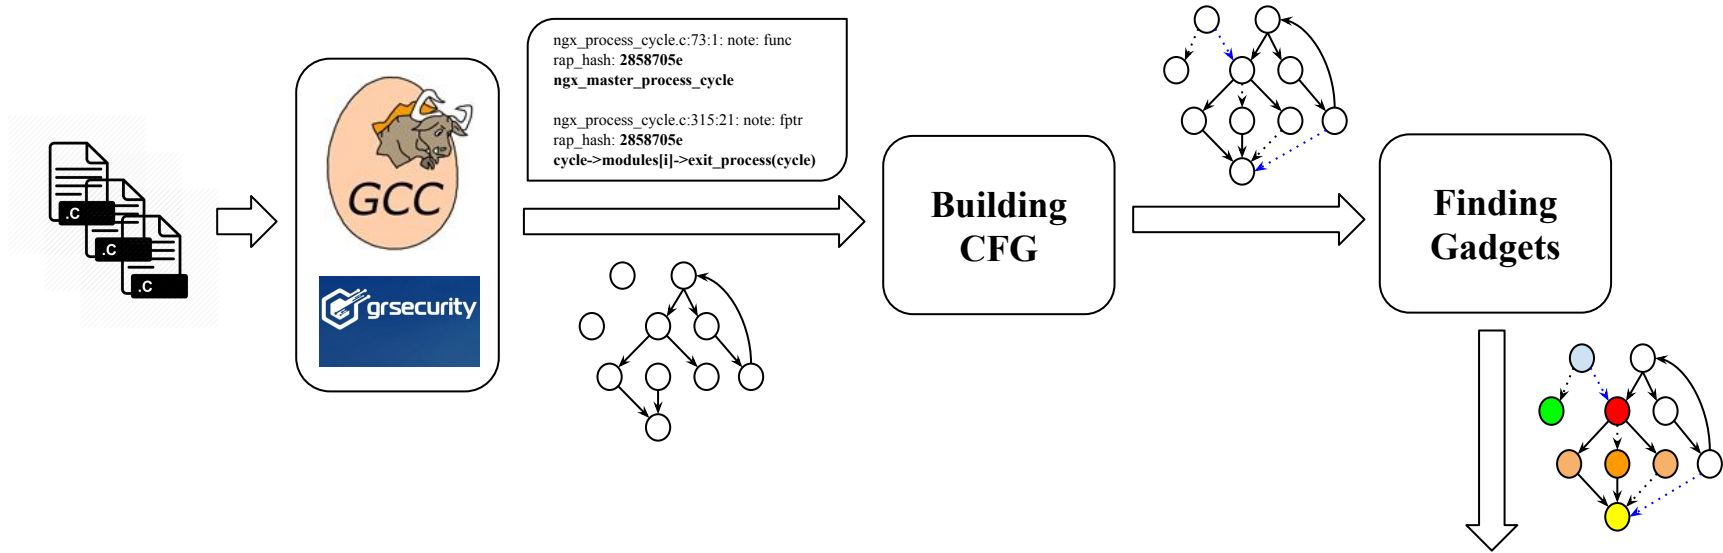


Typed ROP (TROP)



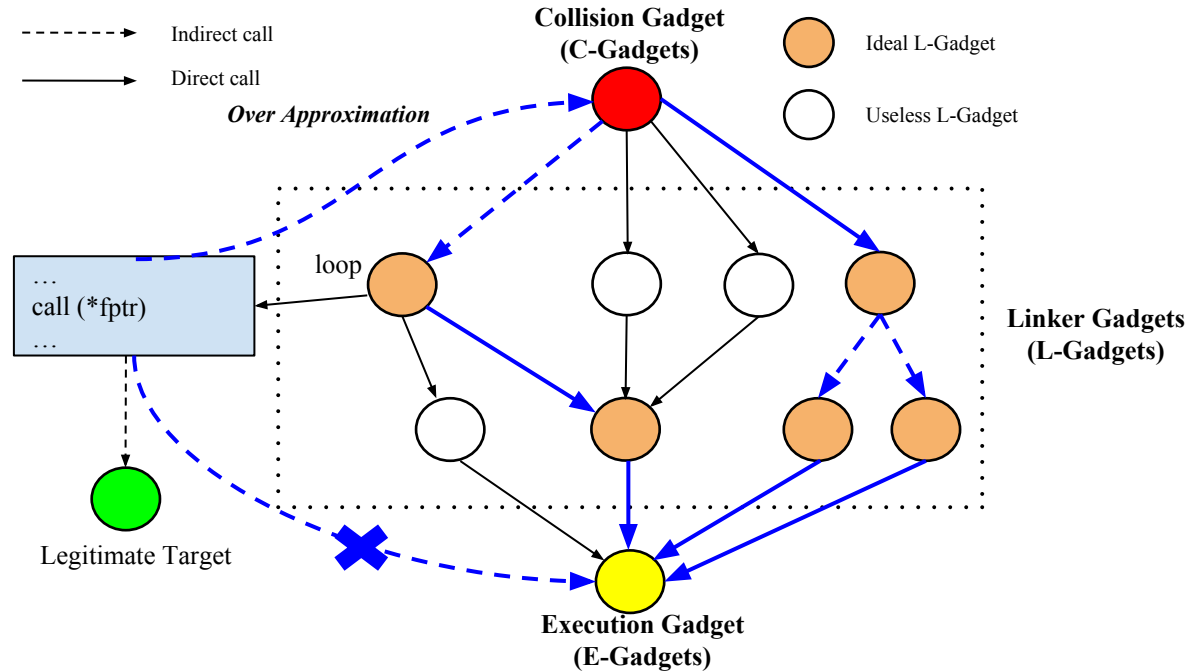


Typed ROP (TROP)



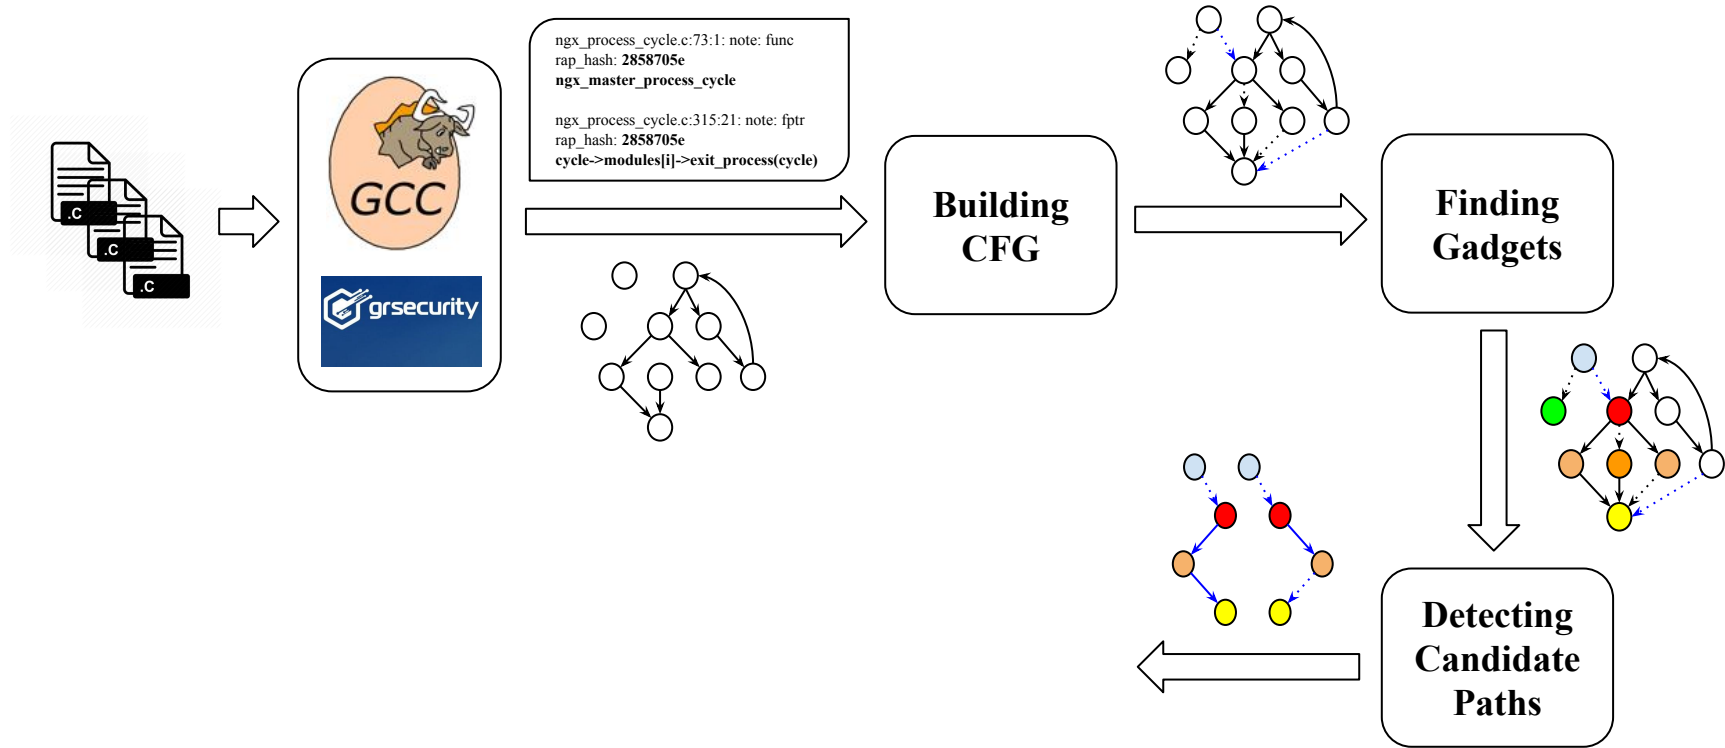


Gadgets



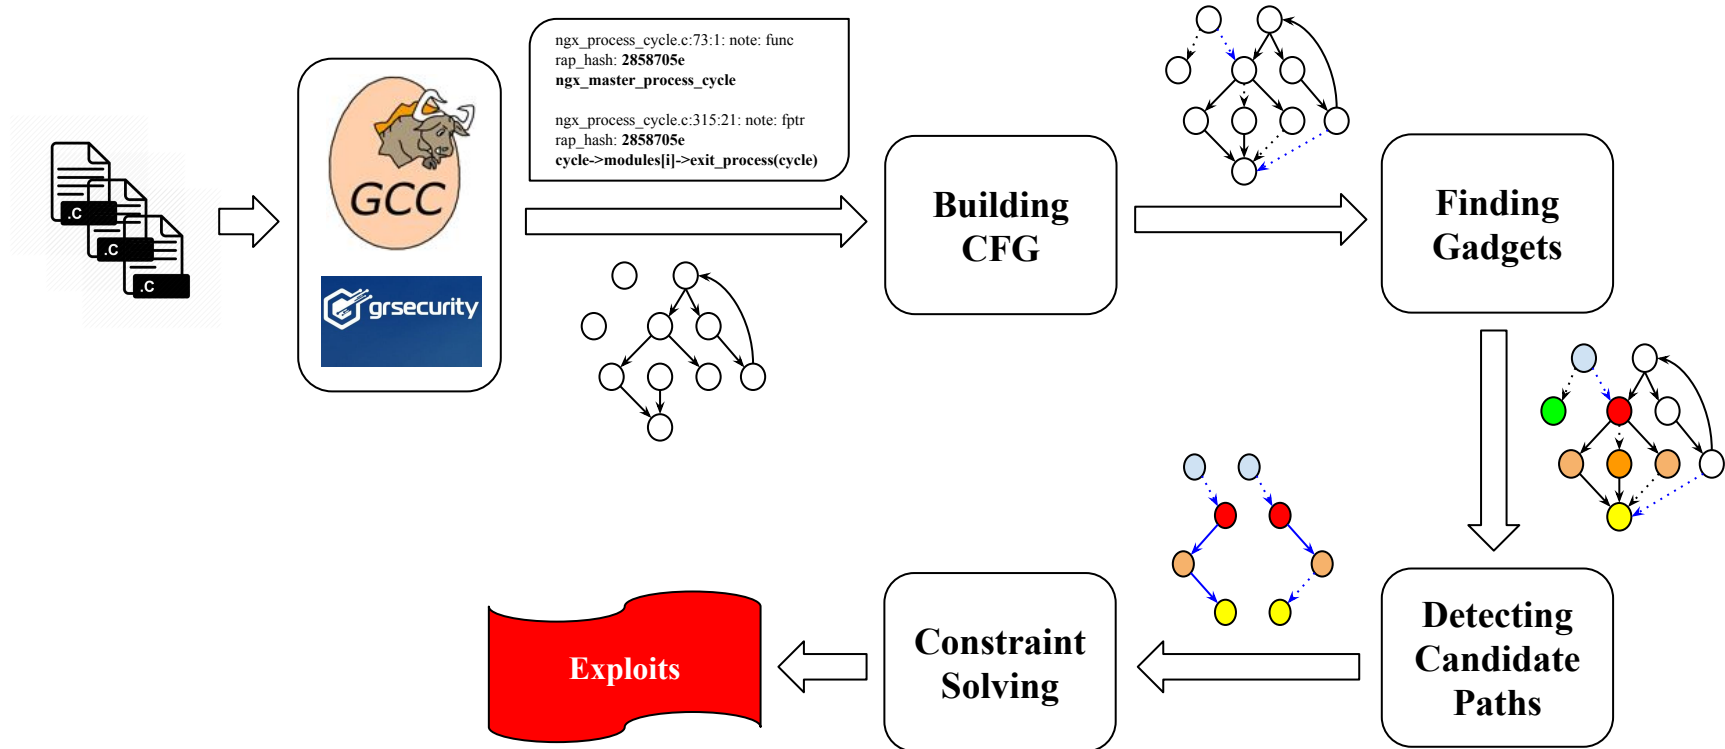


Typed ROP (TROP)



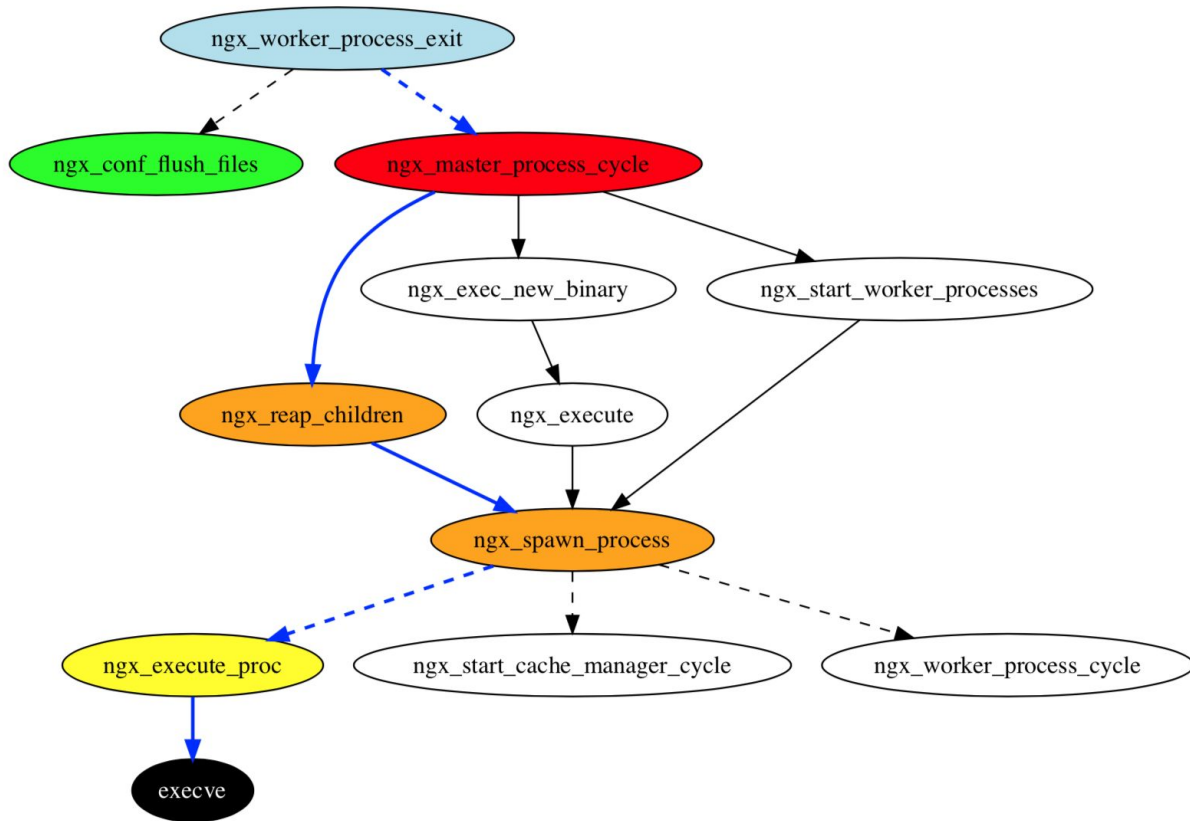


Typed ROP (TROP)





PoC Exploits (Nginx)





PoC Exploits (Nginx)



ngx_worker_process_exit

...
check **2858705e** hash
call (*exit_process)

ngx_conf_flush_files()

hash **2858705e**
...

```
static void
ngx_worker_process_exit(ngx_cycle_t *cycle) {
    ngx_uint_t i;
    ngx_connection_t *c;

    for (i = 0; cycle->modules[i]; i++) {
        if (cycle->modules[i]->exit_process) {
            cycle->modules[i]->exit_process(cycle);
        }
    }
    ...
}
```



PoC Exploits (Nginx)

ngx_worker_process_exit

```
...  
check 2858705e hash  
call (*exit_process)
```

ngx_conf_flush_files()

```
hash 2858705e  
...
```

ngx_master_process_cycle()

```
hash 2858705e  
call ngx_reap_children()  
...
```

C-Gadget

void

```
ngx_master_process_cycle(ngx_cycle_t * cycle) {  
    ...  
    /* By setting this condition to true, the attacker can  
     * reach to the next gadget which is ngx_reap_children()  
     */  
    if (ngx_reap) {  
        ngx_reap = 0;  
        ngx_log_debug0(NGX_LOG_DEBUG_EVENT,  
                        cycle, log, 0, "reap children");  
        live = ngx_reap_children(cycle);  
        ...  
        ...  
    }  
}
```



PoC Exploits (Nginx)

ngx_worker_process_exit

...
check **2858705e** hash
call (*exit_process)

ngx_conf_flush_files()

hash **2858705e**
...

ngx_master_process_cycle()

hash **2858705e**
call ngx_reap_children()
...

C-Gadget

ngx_reap_children()

...
call ngx_spawn_process()
...

L-Gadgets

```
static ngx_uint_t
ngx_reap_children(ngx_cycle_t * cycle) {
    ...
    for (i = 0; i < ngx_last_process; i++) {
        ...
        if (ngx_processes[i].respawn &&
            !ngx_processes[i].exiting &&
            !ngx_terminate &&
            !ngx_quit) {
            if (ngx_spawn_process(cycle,
                ngx_processes[i].proc,
                ngx_processes[i].data,
                ngx_processes[i].name, i)
                == NGX_INVALID_PID) {
                ...
            }
        }
    }
}
```



PoC Exploits (Nginx)

ngx_worker_process_exit

...
check **2858705e** hash
call (*exit_process)

ngx_conf_flush_files()

hash **2858705e**
...

ngx_master_process_cycle()

hash **2858705e**
call ngx_reap_children()
...

C-Gadget

ngx_reap_children()

...
call ngx_spawn_process()
...

L-Gadgets

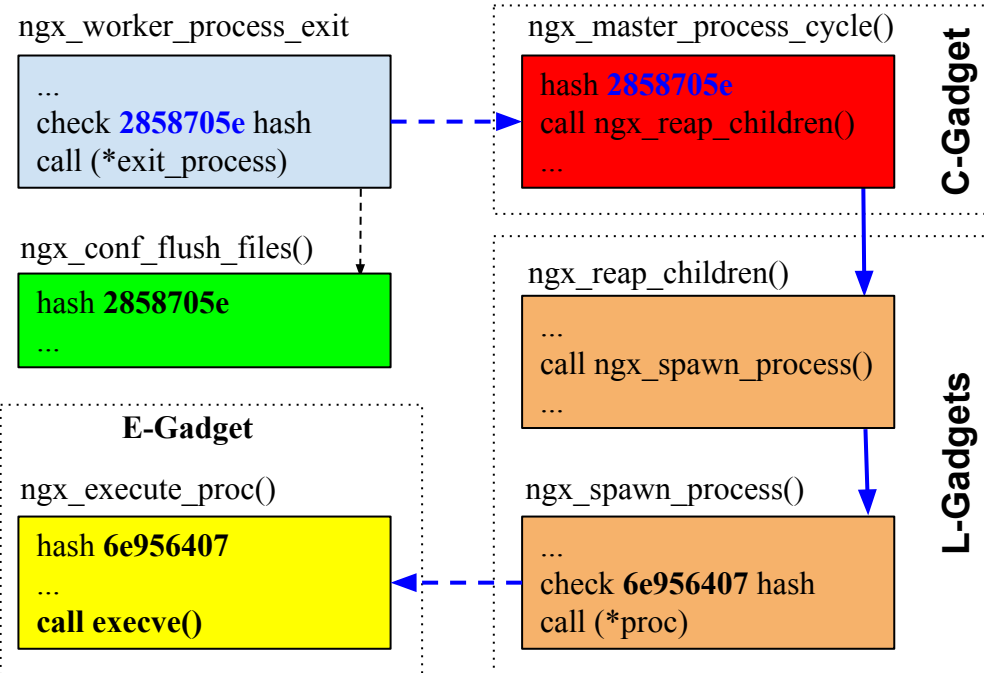
ngx_spawn_process()

...
check **6e956407** hash
call (*proc)

```
ngx_pid_t ngx_spawn_process(ngx_cycle_t *cycle,
    ngx_spawn_proc_pt proc, void *data, char *name,
    ngx_int_t respawn) {
    ...
    switch (pid) {
        case -1:
            ...
        case 0:
            ngx_pid = ngx_getpid();
            proc(cycle, data);
            ...
    }
}
```



PoC Exploits (Nginx)



```
static void
ngx_execute_proc(ngx_cycle_t *cycle, void *data) {
    ngx_exec_ctx_t * ctx = data;

    if (execve(ctx->path, ctx->argv, ctx->envp) == -1) {
        ngx_log_error(NGX_LOG_ALERT,
            cycle -> log,
            ngx_errno,
            "execve() failed while executing %s\ \"%s\"",
            ctx -> name, ctx -> path);
    }

    exit(1);
}
```




Evaluation



- Type Collisions
- Gadget Distribution
- Libc
- Type Checking vs. Points-to Analysis
- Type Diversification
- Practical Challenges



Type Collisions



App	Version	Function Pointer	Call Sites	Functions	Functions w/ Hash	Function Targets		Indirect Calls	
						All	Invalid	All	Invalid
base-passwd	3.5.39	6	6	45	45 (100.0%)	0	0	0	0 (0.0%)
coreutils	8.2	42	80	1,789	682 (38.1%)	116	43	416	110 (26.4%)
e2fsprogs	1.42.13	97	264	1,964	1,243 (63.3%)	251	176	1,383	400 (28.9%)
exim	4.89	43	93	968	607 (62.7%)	88	121	359	165 (46.0%)
findutils	4.6.0	28	52	821	554 (67.5%)	200	89	326	65 (19.9%)
grep	2.25	19	28	460	264 (57.4%)	38	19	113	52 (46.0%)
httpd	2.4.25	248	546	2,800	2,338 (83.5%)	1,332	483	3,915	794 (20.3%)
lighttpd	1.4.45	27	108	899	524 (58.3%)	228	40	830	221 (26.6%)
ncurses	6.0	46	77	1,835	1,045 (56.9%)	156	273	969	397 (41.0%)
nginx	1.10.1	84	290	1,299	977 (75.2%)	610	319	5,977	3,512 (58.8%)
sed	4.2.2	1	1	213	140 (65.7%)	2	0	2	0 (0.0%)
tar	1.28	46	86	1,166	730 (62.6%)	141	166	1,008	754 (74.8%)
util-linux	2.27.1	53	75	3,143	1,681 (53.5%)	211	177	1,060	643 (60.7%)
zlib	1.2.8	5	14	152	108 (71.1%)	5	0	13	0 (0.0%)




Gadget Distribution



App	Version	C-GADGET	L-GADGET	E-GADGET
nginx	1.10.1	8	6	1
httpd	2.4.25	40	19	5
lighttpd	1.4.45	8	29	6
exim	4.90	16	32	7




Libc (Targets)





Libc (Edges)






Type Checking vs. Points-to Analysis

App	Base	Type Checking		Points-to Analysis	
		Total	Invalid	Total	Invalid
base-passwd	0	0	0 (0.0%)	0	0 (0.0%)
coreutils	213	291	78 (26.8%)	308	198 (64.3%)
e2fsprogs	557	861	304 (35.3%)	42	15 (35.7%)
exim	107	212	105 (49.5%)	169	99 (58.6%)
findutils	237	279	42 (15.1%)	448	231 (51.6%)
grep	54	105	51 (48.6%)	108	60 (55.6%)
httpd	2,126	2,870	744 (25.9%)	-	-
lighttpd	327	442	115 (26.0%)	1,096	938 (85.6%)
ncurses	291	558	267 (47.8%)	507	238 (46.9%)
nginx	1,276	2,287	1,011 (44.2%)	-	-
sed	2	2	0 (0.0%)	2	0 (0.0%)
tar	208	664	456 (68.7%)	360	167 (46.4%)
util-linux	311	943	632 (67.0%)	596	465 (78.0%)
zlib	10	10	0 (0.0%)	10	4 (40.0%)




Type Diversification





Type Diversification

- Complicates separate compilation





Practical Challenges



- Mismatch Types
 - **void *** can point to any pointer (e.g., **int ***)
- Support for Assembly Code



Conclusion



- Evaluated RTC from security and practicality perspectives
- Type collisions between function pointers and E-Gadgets are rare
- TROP showed collisions with other functions in a nested fashion can be exploited
- Gadgets for mounting TROP are **abundant** in real-world applications
- RTC is a practical defense but **not sufficient** to prevent control flow hijacking



Questions?



Reza Mirzazade farkhani

reza699@ccs.neu.edu